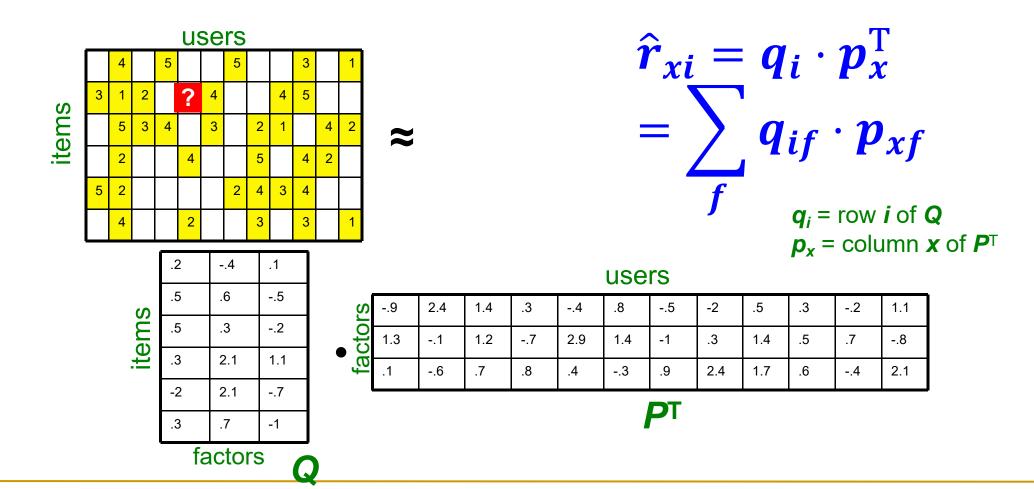
Ratings as Products of Factors



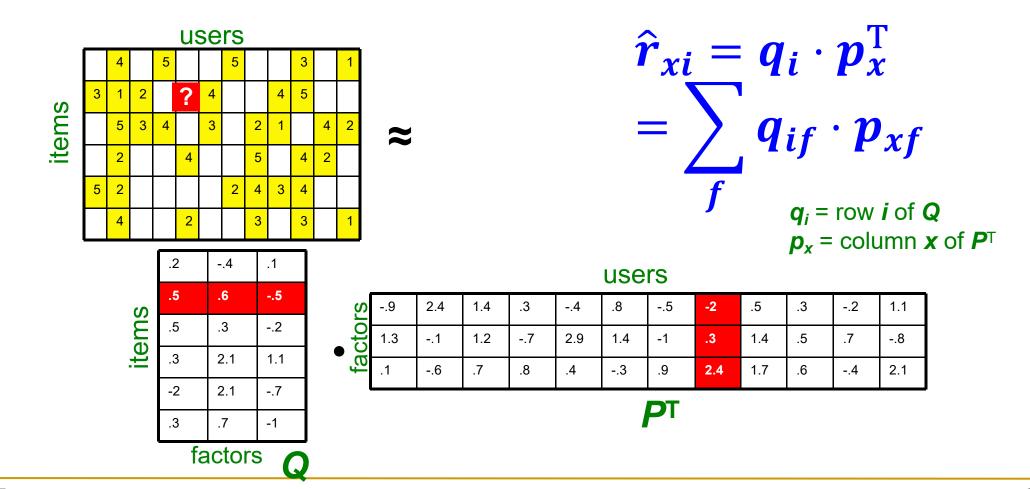
\square How to estimate the missing rating of user x for item i?



Ratings as Products of Factors



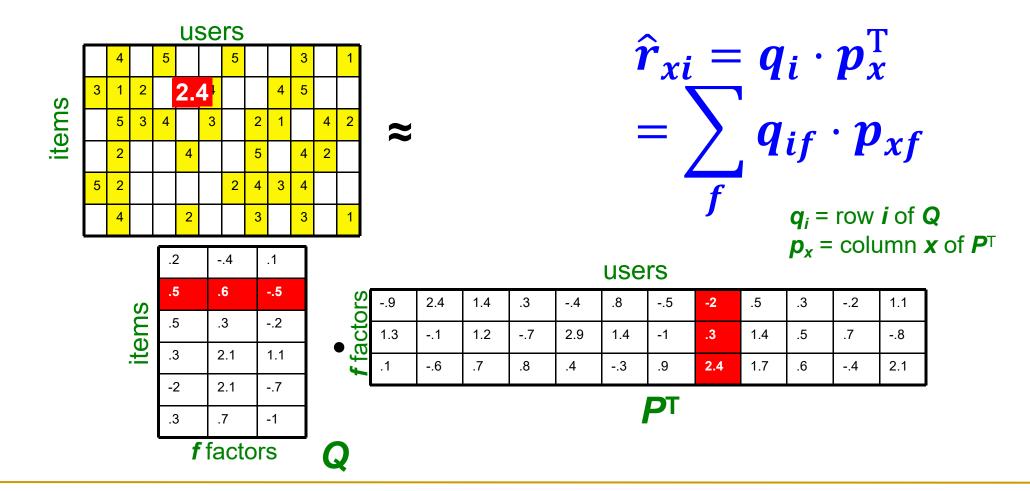
\square How to estimate the missing rating of user x for item i?



Ratings as Products of Factors

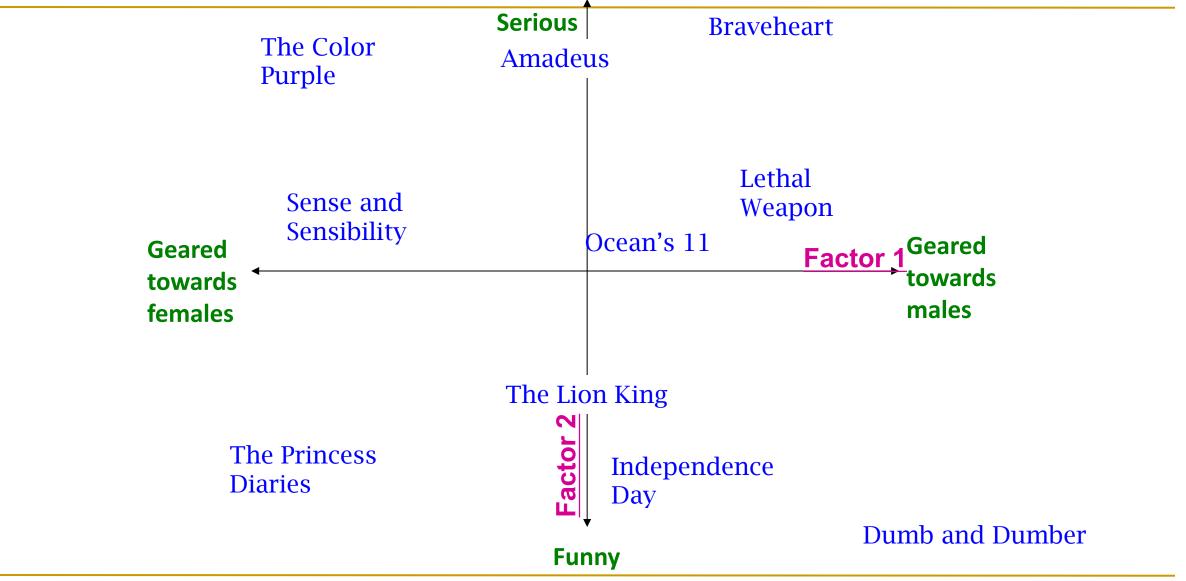


\square How to estimate the missing rating of user x for item i?



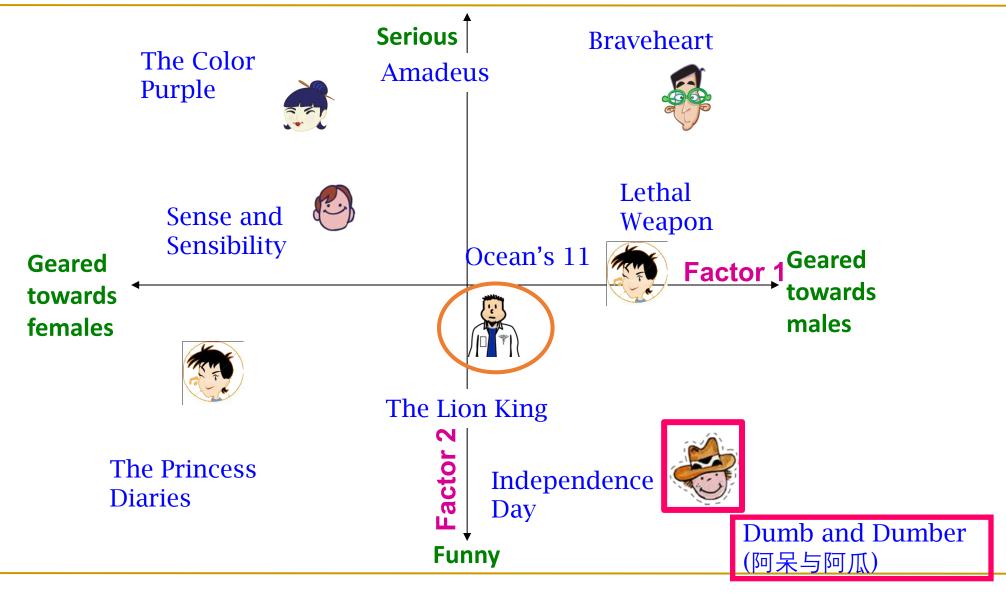
Latent Factor Models





Latent Factor Models





SVD



□SVD (下个章节会详细讲解):

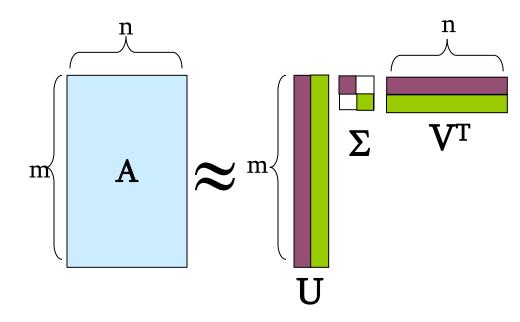
- **▶A**: Input data matrix
- **▶U**: Left singular vecs
- **▶V**: Right singular vecs
- >Σ: Singular values

SVD:
$$A = U \Sigma V^T$$



"SVD" on Netflix data: $R \approx Q \cdot P^T$

$$A = R$$
, $Q = U$, $P^{T} = \Sigma V^{T}$



$$\hat{\boldsymbol{r}}_{xi} = \boldsymbol{q}_i \cdot \boldsymbol{p}_x^T$$

SVD: More good stuff



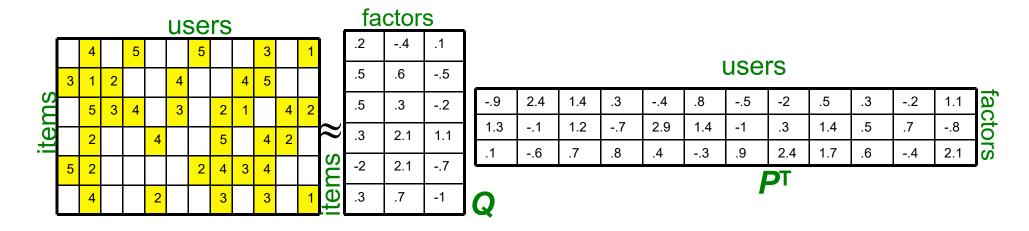
□SVD gives minimum reconstruction error (Sum of Squared Errors, SSE):

$$\min_{U,V,\Sigma} \sum_{ij \in A} \left(A_{ij} - [U\Sigma V^{\mathrm{T}}]_{ij} \right)^{2}$$

- **■Note two things:**
 - >SSE and RMSE are monotonically related:
 - $RMSE = \frac{1}{c}\sqrt{SSE}$ Great news: SVD is minimizing RMSE
 - ➤ Complication: The sum in SVD error term is over all entries (norating in interpreted as zero-rating).
 - But our R has missing entries!

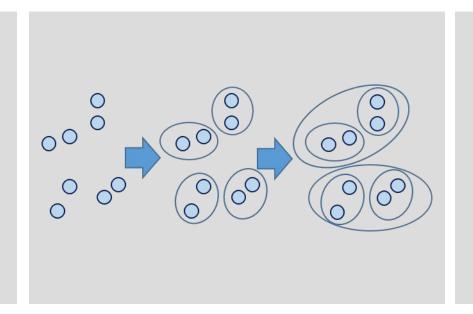
Latent Factor Models





□ Use specialized methods to find *P*, *Q*

- $ightharpoonup \min_{P,O} \sum_{(i,x)\in\mathbb{R}} (r_{xi} q_i \cdot p_x^T)^2$ SSE(平方误差和)
- Note:
 - We don't require cols of P, Q to be orthogonal/unit length
 - *P*, *Q* map users/movies to a latent space



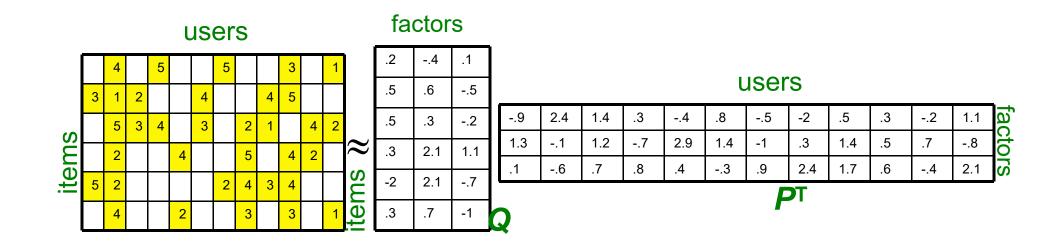
Finding the Latent Factors

Latent Factor Models



□ Our goal is to find P and Q such tat:

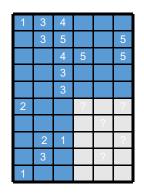
$$\min_{P,Q} \sum_{(i,x)\in R} (r_{xi} - q_i \cdot p_x^T)^2$$



Back to Our Problem



- **■Want to minimize SSE for unseen test data**
- □Idea: Minimize SSE on training data
 - \triangleright Want large k (# of factors) to capture all the signals
 - ▶ But, **SSE** on test data begins to rise for k > 2



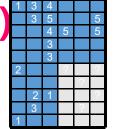
- □This is a classical example of **overfitting(过度拟合, 过拟合):**
 - ➤ With too much freedom (too many free parameters) the model starts fitting noise
 - That is it fits too well the training data and thus not generalizing well to unseen test data

Dealing with Missing Entries



□To solve overfitting we introduce regularization(正则化)

- > Allow rich model where there are sufficient data
- ➤ Shrink aggressively where data are scarce



$$\min_{P,Q} \sum_{training} (r_{xi} - q_i p_x)^2 + \left[\lambda_1 \sum_{x} \|p_x\|^2 + \lambda_2 \sum_{i} \|q_i\|^2 \right]$$
"error"

 λ_1 , λ_2 ... user set regularization parameters(正则化参数); $\|A\|$, Frobenius norm (F范式) $\|A\|_F = \sqrt{\sum\limits_{i=1}^m \sum\limits_{j=1}^n |a_{ij}|^2}$

Note: We do not care about the "raw" value of the objective function, but we care in P,Q that achieve the minimum of the objective

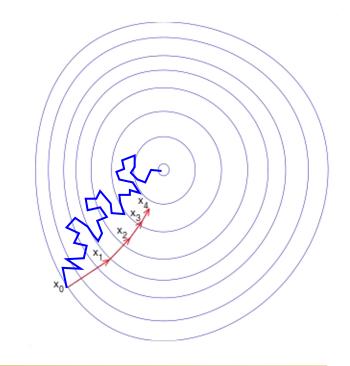
Stochastic Gradient Descent

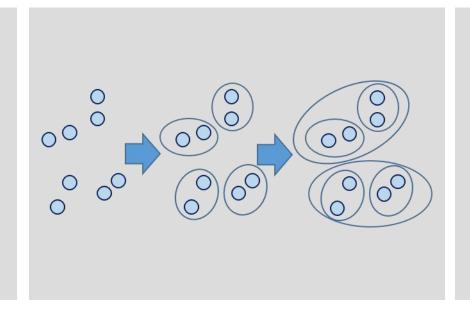


■Want to find matrices P and Q:

$$\min_{P,Q} \sum_{training} (r_{xi} - q_i p_x)^2 + \left[\lambda_1 \sum_{x} ||p_x||^2 + \lambda_2 \sum_{i} ||q_i||^2 \right]$$

- **□**Ans: Stochastic gradient decent(SGD)
 - ➤ SGD, 随机梯度下降, 感兴趣的同学自学





Extending Latent Factor Model to Include Biases

Modeling Biases and Interactions



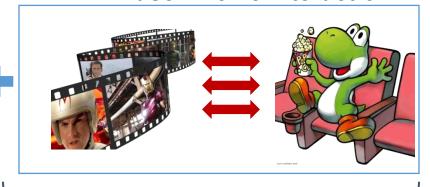
user bias



movie bias



user-movie interaction



Baseline predictor

- Separates users and movies
- Benefits from insights into user's behavior
- Among the main practical contributions of the competition

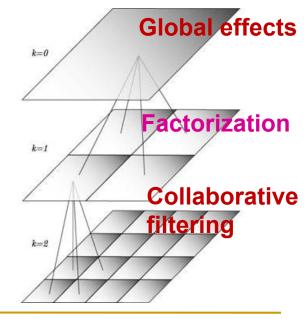
 μ = overall mean rating

 b_x = bias of user x

 b_i = bias of movie i

User-Movie interaction

- Characterizes the matching between users and movies
- Attracts most research in the field
- Benefits from algorithmic and mathematical innovations



Baseline Predictor



 \square We have expectations on the rating by user x of movie i, even without estimating x' s attitude towards movies like i







- Rating scale of user x
- Values of other ratings user gave recently (day-specific mood, anchoring, multi-user accounts)

- (Recent) popularity of movie i
- Selection bias; related to number of ratings user gave on the same day ("frequency")

Putting It All Together



$$r_{xi} = \mu + b_x + b_i + q_i \cdot p_x^{\mathrm{T}}$$

Overall mean rating

Bias for user **x**

Bias for movie *i*

User-Movie interaction

□Example:

- \triangleright Mean rating: $\mu = 3.7$
- You are a critical reviewer: your ratings are 1 star lower than the mean: $b_x = -1$
- ightharpoonupStar Wars (星球大战) gets a mean rating of 0.5 higher than average movie: $b_i = + 0.5$
- ➤ Predicted rating for you on Star Wars:

$$= 3.7 - 1 + 0.5 = 3.2$$

Fitting the New Model



□Solve:

$$\min_{Q,P} \sum_{(x,i)\in R} (r_{xi} - (\mu + b_x + b_i + q_i p_x))^2$$
goodness of fit

$$+ \left(\frac{\lambda_{1}}{1} \sum_{i} \|q_{i}\|^{2} + \lambda_{2} \sum_{x} \|p_{x}\|^{2} + \lambda_{3} \sum_{x} \|b_{x}\|^{2} + \lambda_{4} \sum_{i} \|b_{i}\|^{2} \right)$$
regularization

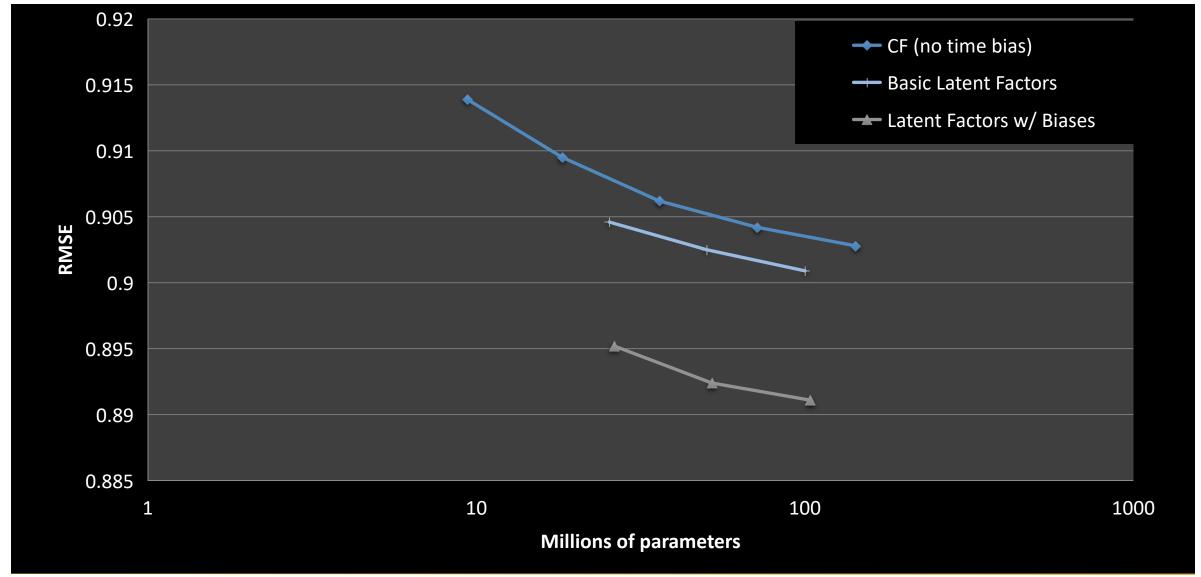
 λ is selected via gridsearch on a validation set

□Stochastic gradient decent (SGD) to find parameters

Note: Both biases b_x , b_i as well as interactions q_i , p_x are treated as parameters (we estimate them)

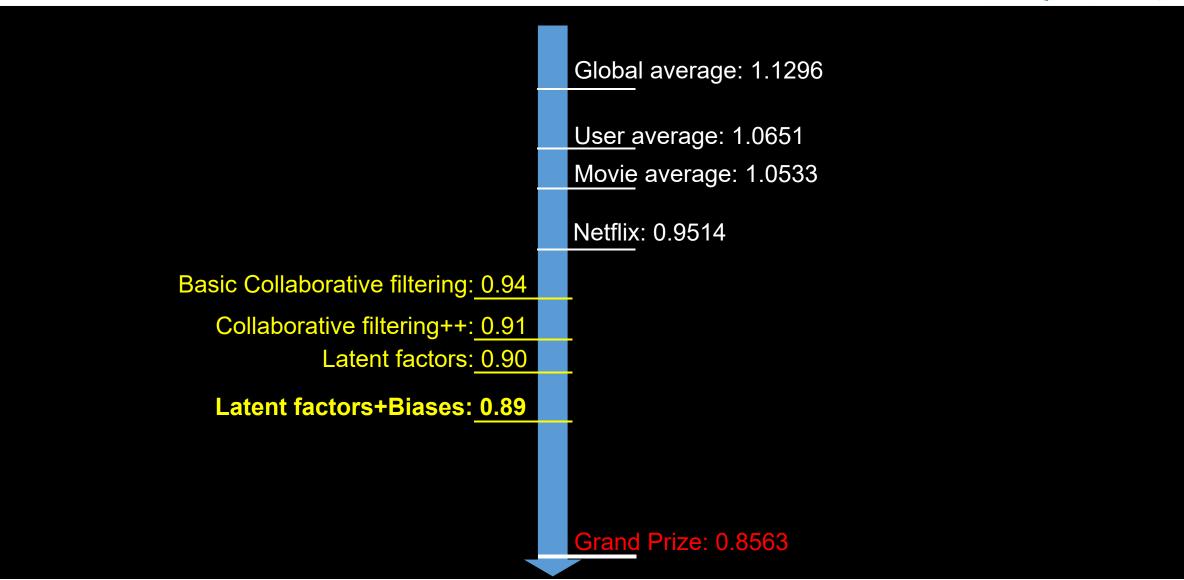
Performance of Various Methods

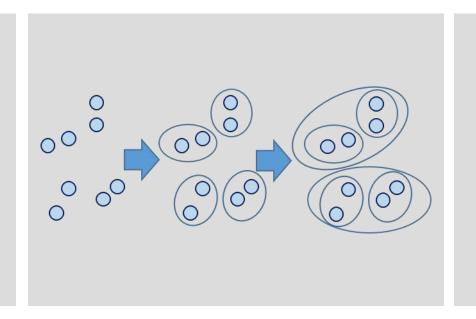




Performance of Various Methods







Add Temporal Biases

Temporal Biases Of Users



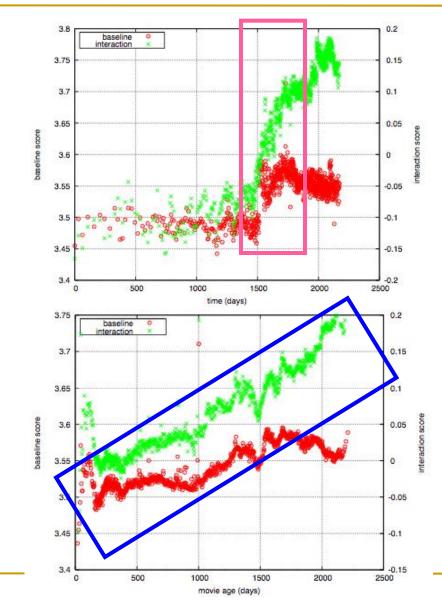
■Sudden rise in the average movie rating (early 2004)

- ➤ Improvements in Netflix
- ➤ GUI improvements
- ➤ Meaning of rating changed

■Movie age

- Users prefer new movies without any reasons
- Older movies are just inherently better than newer ones

Y. Koren, Collaborative filtering with temporal dynamics, KDD '09



Temporal Biases & Factors



□Original model:

$$r_{xi} = \mu + b_x + b_i + q_i \cdot p_x$$

b_x = bias of user x b_i = bias of movie i

■Add time dependence to biases:

$$r_{xi} = \mu + b_x(t) + b_i(t) + q_i \cdot p_x$$

- \triangleright Make parameters b_x and b_i to depend on time
- >(1) Parameterize time-dependence by linear trends
 - (2) Each bin corresponds to 10 consecutive weeks

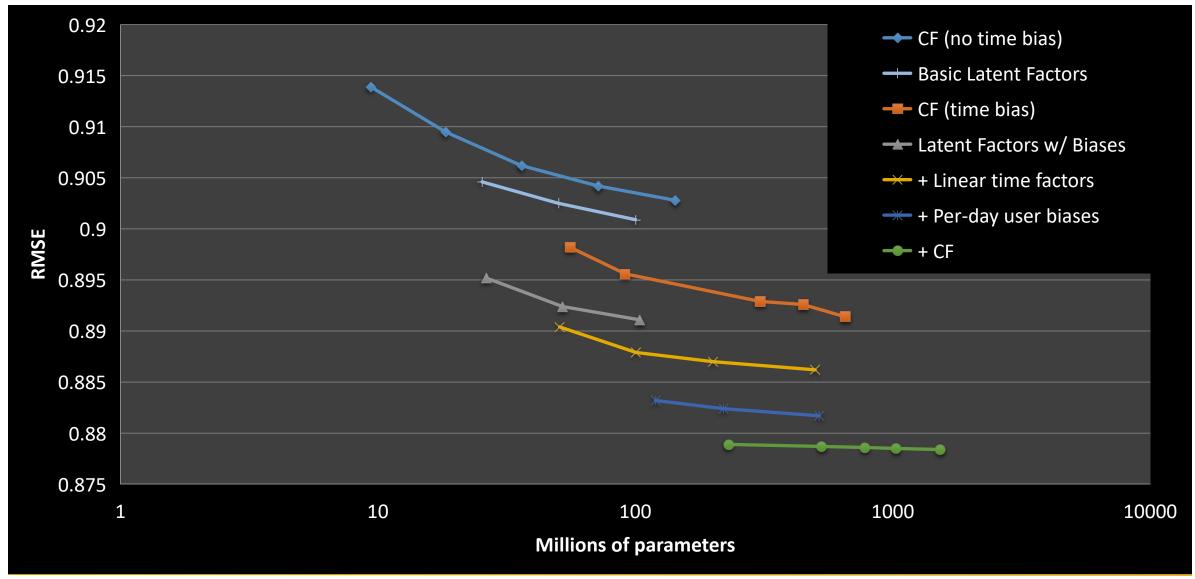
$$b_i(t) = b_i + b_{i,\operatorname{Bin}(t)}$$

□Add temporal dependence to factors

 $p_x(t)$... user preference vector on day t

Adding Temporal Effects





Performance of Various Methods



Global average: 1.1296

User average: 1.0651

Movie average: 1.0533

Netflix: 0.9514

Basic Collaborative filtering: 0.94

Collaborative filtering++: 0.91

Latent factors: 0.90

Latent factors+Biases: 0.89

Latent factors+Biases+Time: 0.876

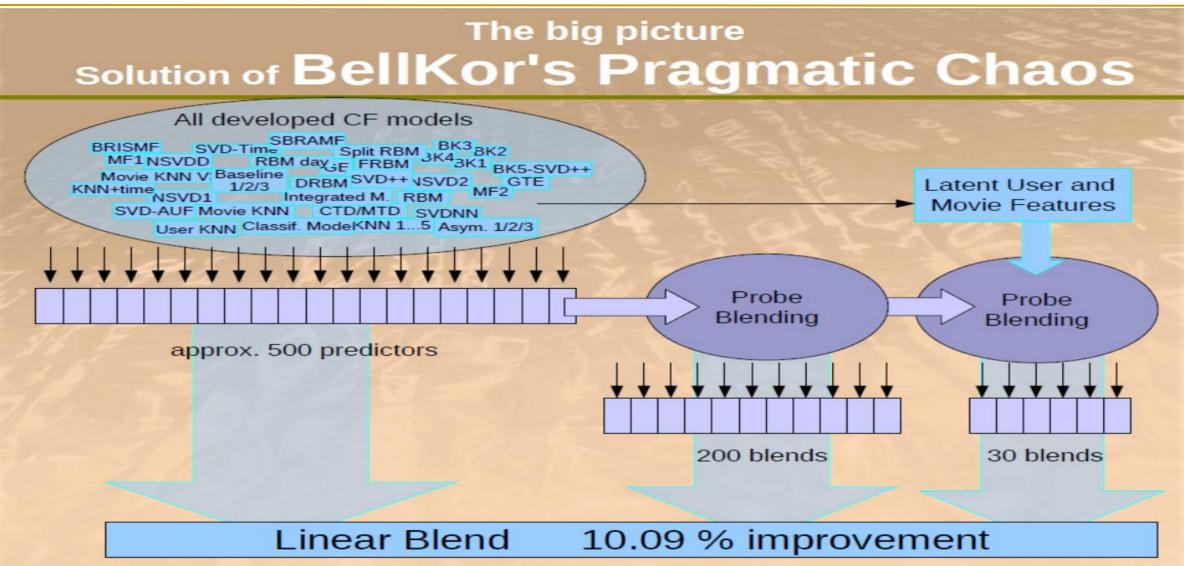
Still no prize!
Getting desperate.

Try a "kitchen sink" approach!

Grand Prize: 0.8563

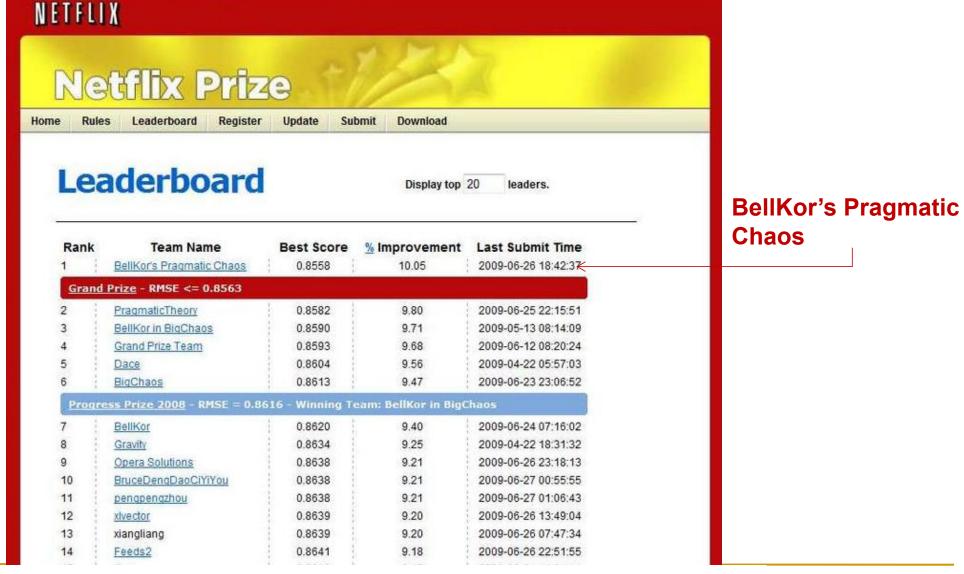
Performance of Various Methods





Standing on June 26th 2009





The Last 30 Days



Ensemble team formed

- ➤ Group of other teams on leaderboard forms a new team
- > Relies on combining their models
- ➤ Quickly also get a qualifying score over 10%

■BellKor

- Continue to get small improvements in their scores
- > Realize that they are in direct competition with Ensemble

■Strategy

- Both teams carefully monitoring the leaderboard
- > Only sure way to check for improvement is to submit a set of predictions
 - This alerts the other team of your latest score

24 Hours from the Deadline



■Submissions limited to 1 a day

➤ Only 1 final submission could be made in the last 24h

■24 hours before deadline...

➤ BellKor team member in Austria notices (by chance) that Ensemble posts a score that is slightly better than BellKor's

☐ Frantic last 24 hours for both teams

- ➤ Much computer time on final optimization
- > Carefully calibrated to end about an hour before deadline

□Final submissions

- > BellKor submits a little early (on purpose), 40 mins before deadline
- > Ensemble submits their final entry 20 mins later
- >....and everyone waits....

Netflix Prize

Home

Rules

Leaderboard

Update

Download

Leaderboard

Showing Test Score. Click here to show quiz score

COMPLETED

Display top 20 \$ leaders.

Rank	Team Name	Best Test Score	% Improvement	Best Submit Time
Grand	Prize - RMSE = 0.8567 - Winning To	sam BellKor's Proof	netic Chees	
1	BellKor's Pragmatic Chaos	0.8567	10.06	2009-07-26 18:18:28
2	The Ensemble	0.8567	10.06	2009-07-26 18:38:22
3	Grand Prize Team	0.0002	9.50	2000 07-10 2 240
4	Opera Solutions and Vandelay United	0.8588	9.84	2009-07-10 01:12:31
5	Vandelay Industries !	0.8591	9.81	2009-07-10 00:32:20
6	PragmaticTheory	0.8594	9.77	2009-06-24 12:06:56
7	BellKor in BigChaos	0.8601	9.70	2009-05-13 08:14:09
8	Dace	0.8612	9.59	2009-07-24 17:18:43
9	Feeds2	0.8622	9.48	2009-07-12 13:11:51
10	BigChaos	0.8623	9.47	2009-04-07 12:33:59
11	Opera Solutions	0.8623	9.47	2009-07-24 00:34:07
12	BellKor	0.8624	9.46	2009-07-26 17:19:11
Progr	ess Prize 2008 - RMSE = 0.8627 - W	inning Team: BellKo	r in BigChaos	
13	xiangliang	0.8642	9.27	2009-07-15 14:53:22
14	Gravity	0.8643	9.26	2009-04-22 18:31:32
15	Ces	0.8651	9.18	2009-06-21 19:24:53
16	Invisible Ideas	0.8653	9.15	2009-07-15 15:53:04
17	Just a guy in a garage	0.8662	9.06	2009-05-24 10:02:54
18	J Dennis Su	0.8666	9.02	2009-03-07 17:16:17
19	Craig Carmichael	0.8666	9.02	2009-07-25 16:00:54
20	acmehill	0.8668	9.00	2009-03-21 16:20:50

Million \$ Awarded Sept 21st 2009





Some slides and plots borrowed from Yehuda Koren, Robert Bell and Padhraic Smyth. **Further reading:** Y. Koren, Collaborative filtering with temporal dynamics, KDD '09

http://www2.research.att.com/~volinsky/netflix/bpc.html

http://www.the-ensemble.com/

Chapter 6 总结



- **□**Three Key Problems for Recommender systems:
- □(1) Gathering "known" ratings for matrix
- □(2) Extrapolate unknown ratings from the known ones
 - **▶1)** Content-based
 - >2) Collaborative-based
 - **▶3)** Latent factor based
- **□(3) Evaluating extrapolation methods**