

明 求是创新 厚学 新

PARALLEL AND SEQUENTIAL ALGORITHMS AND DATA STRUCTURES

LECTURE 2 SPARC and examples



华中科技大学
HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

SYNOPSIS

- Functional Algorithms
- The Lambda Calculus
- The SPARC Language
 - Syntax and Semantics
 - Type System of SPARC
- Example
 - Threads, Concurrency, and Parallelism
 - Critical Sections and Mutual Exclusion



华中科技大学
HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Introduction

- We define algorithms and data structures using **nested parallelism** in conjunction with a **functional programming style**
 - This is the best way to capture the core ideas of algorithms and **parallelism in a concise, clear, safe, and precise way**
 - will be useful in a broad set of programming languages



华中科技大学
HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Nested parallelism

- Nested parallelism (or nested fork-join parallelism)
 - a style of parallelism in which any task can **fork a set of new child tasks** to run in parallel
 - When forking, the parent task suspends
 - when all the child tasks finish, they “join”, and the parent continues
- **sufficiently powerful** to capture the parallelism in most of the algorithms needed for the purpose of this book
- For dynamic programming, we diverge slightly from this model to a somewhat more general model



Functional Algorithm

- a style of programming in which functions act **like mathematical functions** (a mapping from domain to a codomain, and no side effects), and can be used as values (can be passed around, stored as data, and created on the fly)
 - have **no side effects**, parallelism is inherently safe and deterministic
 - allows **powerful abstractions**



华中科技大学
HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

德

- We use a minimal, perhaps “toy”, language called SPARC to describe algorithms and data structures.
 - SPARC has structures for supporting **nested parallelism** and supports only **functional programming**
 - is an extension of the **lambda-calculus**



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Function vs. Algorithm

- Functions are more accurately algorithms
 - mapping from inputs to outputs (**mathematical functions**)
 - specify the mechanism (code) by which the output is generated from the input



华中科技大学
HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

The Lambda Calculus

- developed *by Alonzo Church* in the early 30s, is arguably the first general purpose “programming language”
- Is a pure language, only supporting pure functions, and it fully supports higher-order functions
 - it is very simple with only **three types of expressions**, and **one rule for “evaluation”**
- It captures many of the core ideas of modern programming languages
- **SPARC**, is effectively an extended and typed lambda calculus



Syntax and Semantics

- Definition (Syntax of the Lambda Calculus)
 - The lambda calculus consists of expressions e
 - a **variable**, such as x, y, z, \dots ,
 - a **lambda abstraction**, written as $(\lambda x. e)$, where x is a variable name and e is an expression, or
 - an **application**, written as $(e_1 e_2)$, where e_1 and e_2 are expressions



Syntax and Semantics

- **Definition (Beta Reduction)**

- For any application for which the left hand expression is a lambda abstraction, **beta reduction** “applies the function” by making the transformation:
- $(\lambda x . e_1) e_2 \rightarrow e_1[x/e_2]$
- where $e_1[x/e_2]$ roughly means for every (free) occurrence of x in e_1 , substitute it with e_2



明

德

學



Church-Turing Hypothesis

- Story
 - In the early 30s, soon after he developed the language, Church argued that anything that can be “effectively computed” can be computed with the lambda calculus, and therefore it is a universal mechanism for computation
 - A few years later, Alan Turing developed the Turing machine and showed its equivalence to the lambda calculus that the concept of universality became widely accepted
- Church-Turing hypothesis: anything that can be computed can be computed with the lambda calculus, or equivalently the Turing machine
- Church-Turing complete: any computational model that is computationally equivalent to the lambda calculus



華中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Parallelism and Reduction Order

- the lambda calculus is inherently parallel
- Definition (Call-by-Value)
 - beta reduction is only applied to $(\lambda x . e_1) e_2$ if the expression e_2 is a value, i.e., e_2 is evaluated to a value (lambda abstraction) first, and then beta reduction is applied
- Definition (Call-by-Need)
 - beta reduction is applied to $(\lambda x . e_1) e_2$ even if e_2 is not a value (it could be another application)
 - If during beta reduction e_2 is copied into each variable x in the body, this reduction order is called call-by-name , and if e_2 is shared, it is called call-by-need



华中科技大学
HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Parallelism and Reduction Order

- **Call-by-value is an inherently parallel reduction order**
 - This is because in an expression $(e_1 e_2)$ the two subexpressions can be evaluated (reduced) in parallel, and when both are fully reduced we can apply beta reduction to the results
- **call-by-need is inherently sequential**
 - In an expression $(e_1 e_2)$ only the first subexpression can be evaluated and when completed we can apply beta reduction to the resulting lambda abstraction by substituting in the second expression
 - Therefore the second expression cannot be evaluated until the first is done without potentially changing which reductions are applied



SPARC

- SPARC is a **parallel** and “strict” functional language used throughout the book for specifying algorithms
- is formed from a mixture of several languages
 - ML class of languages : SML, Caml, F#
 - Mathematical notation
 - English descriptions



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

The SPARC Language: Syntax and Semantics

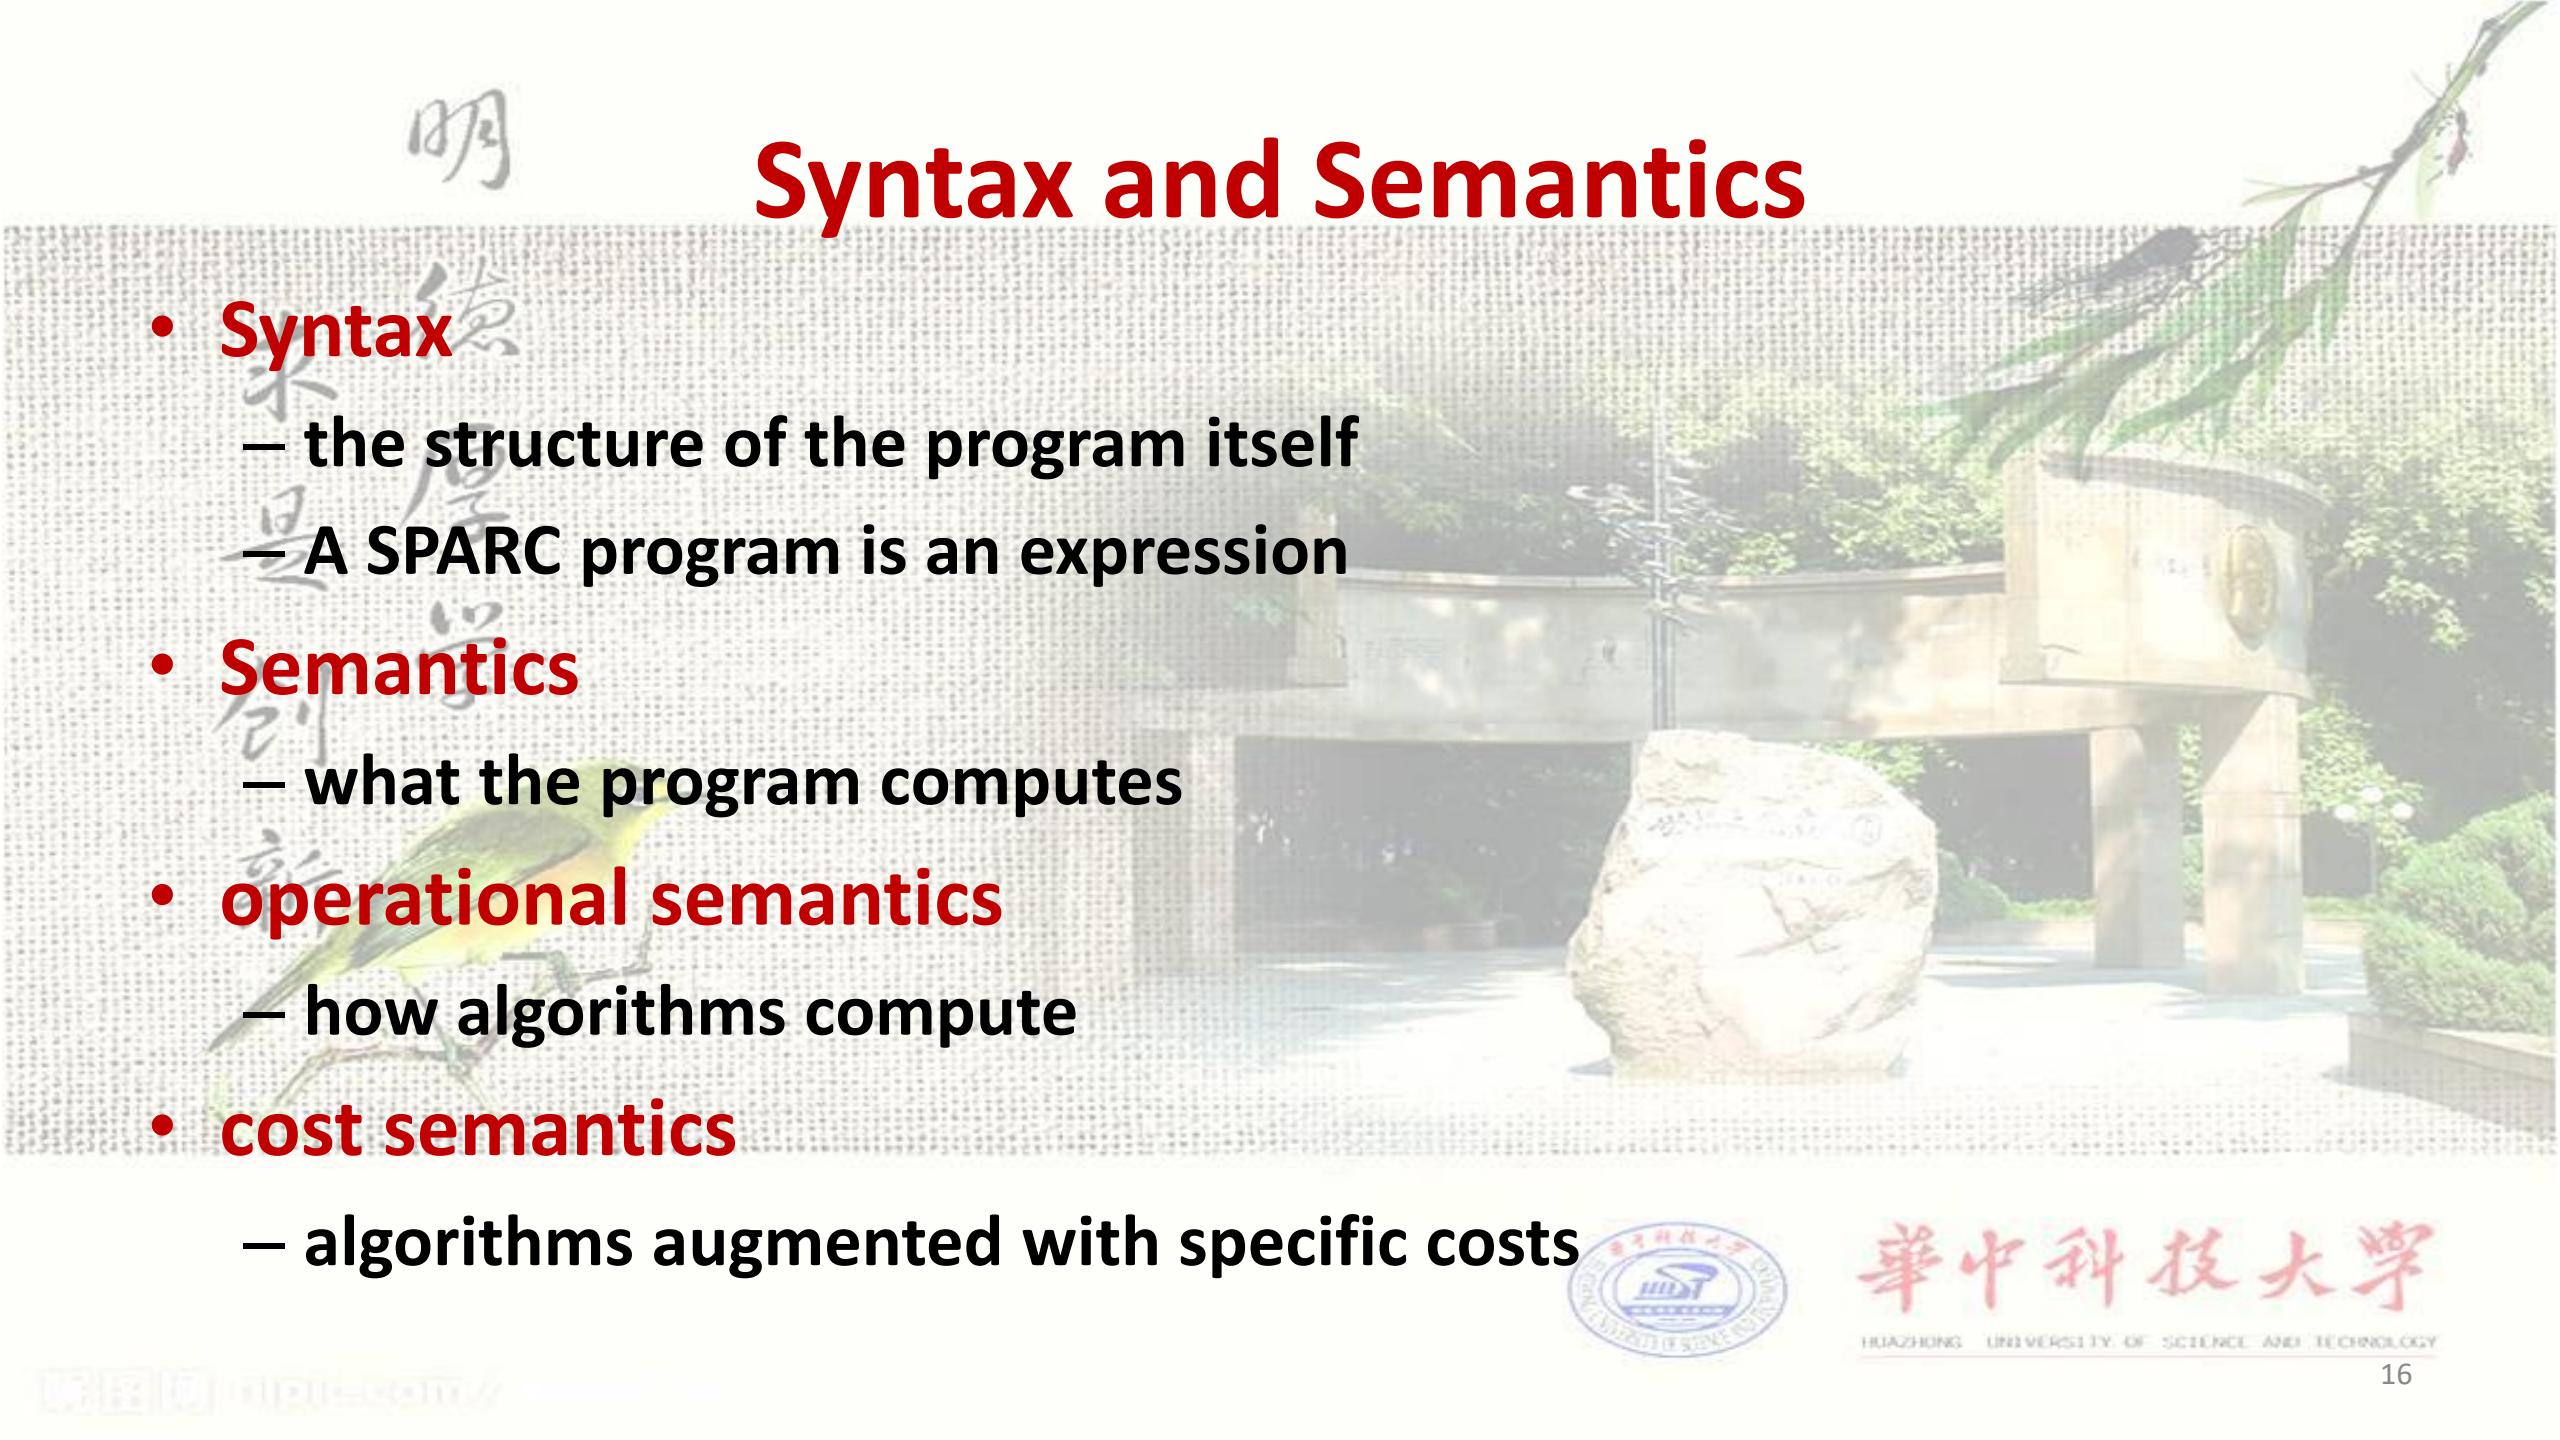
- How to understand?

```
1  listofOne () =
2  let
3      x = 1
4  type t = Int
5  type list = Nil | Cons of (t * list)
6  in
7      Cons(x,Nil)
8  end
```



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY



Syntax and Semantics

- **Syntax**
 - the structure of the program itself
 - A SPARC program is an expression
- **Semantics**
 - what the program computes
- **operational semantics**
 - how algorithms compute
- **cost semantics**
 - algorithms augmented with specific costs



Syntax Definition (SPARC expressions)

Identifier	$id ::= \dots$	
Variables	$x ::= id$	
Type Constructors	$tycon ::= id$	
Data Constructors	$dcon ::= id$	
Patterns	$p ::= x$ (p) p_1, p_2 $dcon (p)$	variable parenthesis pair data pattern
Types	$\tau ::= \mathbb{Z}$ \mathbb{B} $\tau^{[*\tau]^+}$ $\tau \rightarrow \tau$ $tycon$ dty	integers booleans products functions type constructors data types
Data Types	$dty ::= dcon [of \tau]$ $dcon [of \tau] \mid dty$	
Operations	$op ::= + \mid - \mid * \mid - \dots$	
Bindings	$b ::= x(p) = e$ $p = e$ $\text{type } tycon = \tau$ $\text{type } tycon = dty$	bind function bind pattern bind type bind datatype

Values	$v ::= 0 \mid 1 \mid \dots$ $\mid -1 \mid -2 \mid \dots$ $\mid \text{true} \mid \text{false}$ $\mid \text{not} \mid \dots$ $\mid \text{and} \mid \text{plus} \mid \dots$ $\mid v_1, v_2$ $\mid (v)$ $\mid dcon (v)$ $\mid \lambda p . e$	integers integers booleans unary operations binary operations pairs parenthesis constructed data lambda functions
Expression	$e ::= x$ $\mid v$ $\mid e_1 \text{ op } e_2$ $\mid e_1, e_2$ $\mid e_1 \parallel e_2$ $\mid (e)$ $\mid \text{case } e_1 [\mid p \Rightarrow e_2]^+$ $\mid \text{if } e_1 \text{ then } e_2 \text{ else } e_3$ $\mid e_1 \ e_2$ $\mid \text{let } b^+ \text{ in } e \text{ end}$	variables values infix operations sequential pair parallel pair parenthesis case conditionals let expressions

Identifier

- In SPARC, **variables**, **type constructors**, and **data constructors** are given a name, or an **identifier**
 - An identifier consist of only alphabetic and numeric characters (a-z, A-Z, 0-9), the underscore character (“_”), and optionally end with some number of “primes”
 - E.g.: x', x1, xl, myVar, myType, myData, and my_data

<i>Identifier</i>	id	$::=$	\dots
<i>Variables</i>	x	$::=$	id
<i>Type Constructors</i>	$tycon$	$::=$	id
<i>Data Constructors</i>	$dcon$	$::=$	id



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Patterns

- **variables and data constructors can be used to construct more complex patterns over data**
 - For example, a pattern can consist of a **pair (x, y)** or a **triple of variables (x, y, z)**, or it can consist of a data constructor **Cons(x)** or **Cons(x,y)**

Patterns

$$\begin{array}{lcl} p & := & x \\ & | & (p) \\ & | & p_1, p_2 \\ & | & dcon(p) \end{array}$$

variable

parenthesization

pair

data pattern



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Built-in Types

- Types of SPARC include **base types** such as **integers Z**, **booleans B**, **product types** such as $\tau_1 * \tau_2 ... \tau_n$, **function types** $\tau_1 \rightarrow \tau_2$ with domain τ_1 and range τ_2 , as well as **user defined data types**

Types	τ	$\vdash \tau$	
	$= \mathbb{Z}$		integers
	$ \quad B$		booleans
	$ \quad \tau [*\tau]^+$		products
	$ \quad \tau \rightarrow \tau$		functions
	$ \quad tycon$		type constructors
	$ \quad dty$		data types



Data Types

- In addition to built-in types, a program can define *new data types* as a union of tagged types, also called variants, by “unioning” them via distinct data constructors
 - the following data type defines a point as a two-dimensional or a three-dimensional coordinate of integers

```
type point = PointTwo of Z * Z  
           | Point3D of Z * Z * Z
```



Recursive Data Type

- In SPARC *recursive data types* are relatively easy to define and compute with
 - we can define a point list data type as follows

```
type plist = Nil | Cons of point * plist.
```

- Based on this definition the list defines a list consisting of three points

```
Cons(PointTwo(0, 0),  
      Cons(PointTwo(0, 1),  
            Cons(PointTwo(0, 2), Nil)))
```



Option Type

- *Option types* for natural numbers can be defined as follows

```
type option = None | Some of N
```

- option types for integers

```
type intOption = INone | ISome of Z
```



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Values

- The irreducible units of computation include **natural numbers**, **integers**, **Boolean values**, **unary primitive operations**, **binary operations**, **constant-length tuples**, **data constructors applied to values**

Values	v	$\vdash v = 0 \mid 1 \mid \dots$	integers
		$\mid -1 \mid -2 \mid \dots$	integers
		$\mid \text{true} \mid \text{false}$	booleans
		$\mid \text{not} \mid \dots$	unary operations
		$\mid \text{and} \mid \text{plus} \mid \dots$	binary operations
		$\mid v_1, v_2$	pairs
		$\mid (v)$	parenthesis
		$\mid dcon\ (v)$	constructed data
		$\mid \text{lambda}\ p.\ e$	lambda functions



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Values

- As a functional language, SPARC treats *all function as values*
 - The anonymous function *lambda p. e*
 - e is a function whose arguments are specified by the pattern p, and whose body is the expression e
- The function *lambda x. x + 1* takes a single variable as an argument and adds one to it.
- The function *lambda (x, y). x* takes a pairs as an argument and returns the first component of the pair.



Expressions

- Expressions, denoted by e and **variants** (with subscript, superscript, prime), are defined inductively
- **infix expression : $e_1 \ op \ e_2$**
 - Op include + (plus), – (minus), * (multiply), / (divide), < (less), > (greater), or, and and
 - $f(e_1, e_2)$
- We use standard **precedence rules** on the operators to indicate their parsing
 - $3 + 4 * 5 = 3 + (4 * 5)$
 - all operators are left associative unless stated otherwise
 - $5 - 4 + 2$ evaluates to $(5 - 4) + 2 = 3$ not $5 - (4 + 2) = -1$



Sequential and Parallel Composition

- two special infix operators: “,” and “||”

— **comma operator or sequential composition**

- $(e1, e2)$

lambda $(x, y).$ $(x * x, y * y)$

- evaluates $e1$ and $e2$ sequentially, one after the other, returns the ordered pair consisting of the two resulting values

— **parallel operator or parallel**

- $(e1 || e2)$

lambda $(x, y).$ $(x * x || y * y)$

- evaluates $e1$ and $e2$ in parallel, at the same time, and returns the ordered pair consisting of the two resulting values

— **The two operators are identical in terms of their return values**



Case Expressions and Conditionals

- A **case expression** such as

```
case e1
| Nil ⇒ e2
| Cons (x, y) ⇒ e3
```

- **Conditionals**
 - *if-then-else expression*



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

才 德 厚

- A **function application** , $e_1 e_2$,
 - applies the function generated by evaluating e_1 to the value generated by evaluating e_2

(lambda (x, y). x/y) (8, 2)

(lambda (f, x). $f(x, x)$) (plus, 3)

(lambda x. (lambda y. $x + y$)) 3

= ?



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明德厚學
知行合一

Bindings

- The *let expression*

```
let b+in e end,
```

- *Variable binding*

- A *function binding* , $x(p)=e$

- Each *type binding* equates a type to a base type or a data type.

```
1 let
2   x = 2 + 3
3   f(w) = (w × 4, w - 2)
4   (y, z) = f(x - 1)
5 in
6   x + y + z
7 end
```

```
1 let
2   f(i) = if (i < 2) then i else i × f(i - 1)
3 in
4   f(5)
5 end
```



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Binding

let

```
type point = PointTwo of Z * Z  
| PointThree of Z * Z * Z
```

```
injectThree (PointTwo (x, y)) = PointThree (x, y, 0)
```

```
projectTwo (PointThree (x, y, z)) = PointTwo (x, y)
```

```
compose f g = f g
```

```
p0 = PointTwo (0, 0)
```

```
q0 = injectThree p0
```

```
p1 = (compose projectTwo injectThree) p0
```

in

```
(p0, q0)
```

end

```
type tree = Leaf of Z | Node of (tree, Z, tree)
```

```
find (t, x) =
```

```
case t
```

```
| Leaf y ⇒ x = y
```

```
| Node (left, y, right) ⇒
```

```
if x = y then
```

```
return true
```

```
else if x < y then
```

```
find (left, x)
```

```
else
```

```
find (right, x)
```



Binding

- Syntactic sugar for loops, recursion and while

```
sumSquares(n) =  
  let  
    (n, s) =  
      start (n, 0) and  
      while n > 0 do  
        s = s + n2  
        n = n - 1  
      in s end
```



```
sumSquares(n) =  
  let  
    f(n, s) = if not (n > 0) then (n, s)  
              else let  
                      s = s + n × n  
                      n = n - 1  
                    in f(n, s) end  
    (n, s) = f(n, 0)  
  in s end
```

明

Binding

Example 1.43. The following is a recursive variant type defining binary trees.

```
type tree = Leaf | Node of tree * int * tree
```

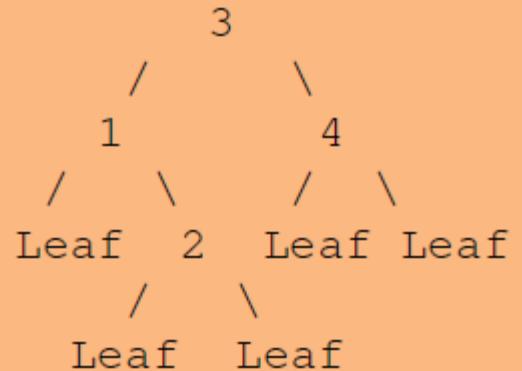
We can then create a tree using, for example

```
Node(Node(Leaf, 1, Node(Leaf, 2, Leaf)), 3, Node(Leaf, 4, Leaf))
```

How many notes?

```
countLeaves(T) =  
  case T of  
    Leaf => 1  
  | Node of (L, R) => countLeaves(L) + countLeaves(R)
```

which corresponds to the tree:



counts the number of leaves of a tree recursively

SYNOPSIS

- Functional Algorithms
- The Lambda Calculus
- The SPARC Language
 - Syntax and Semantics
 - Type System of SPARC
- Example
 - Threads, Concurrency, and Parallelism
 - Critical Sections and Mutual Exclusion



华中科技大学
HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Threads

- A **thread**, short for **thread of execution**: a computation that executes a given piece of code.
- ***multithreaded*** : A program that uses multiple threads
- **two operations on threads**
 - The operation ***spawn*** takes an expression, creates a thread to execute that expression, and returns the thread
 - Once spawned the thread starts executing concurrently with other threads in the program
 - The operation ***sync*** takes a thread and waits until that thread completes its execution



明

Threads

```
let t = spawn (lambda () . fib n )
      u = spawn (lambda () . fib 2n )
      (( ), ( )) = (sync t, sync u )
in () end
```

```
fib x =
  if x ≤ 1 then x
  else fib (x - 1) + fib (x - 2)
```

- Any question? has no way of communicating the result back
 - The sync operations ensure that the results are computed by waiting for the threads to complete

```
let (r, s) = (ref 0, ref 0)
      t = spawn (lambda () . r ← fib n )
      u = spawn (lambda () . s ← fib 2n )
      (( ), ( )) = (sync t, sync u )
in (!r, !s)
end
```



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Thread Scheduler

- Multithreaded programs rely on a *thread scheduler* or scheduler for short, to execute the spawned threads to completion
 - At a given time, a scheduler can execute any subset of the spawned threads that are ready to execute, i.e., they are not waiting other threads



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Concurrency and Parallelism

- ***concurrency problem*** if its specification involves multiple things happening at the same ***item***
 - ***Multithreading*** is usually the technique of choice for solving concurrency
 - Using a ***thread scheduler***, such multithreaded implementations can be made to work on sequential hardware, such as a computer with any number of processors



Parallelism

- An algorithm is **parallel** if it performs multiple tasks **at the same time**
 - Both concurrent and non-concurrent problems typically accept parallel algorithms (solutions)
- **Concurrency** versus **Parallelism**
 - concurrency is a property of a “problem”
 - parallelism is a property of an implementation or a “solution”



明

德

道

Parallel Fibonacci

- The function ***fib*** below computes the ***xth*** Fibonacci number in parallel

```
fib x dest =
  if x ≤ 1 then
    dest ← x
  else
    let
      (da, db) = (ref 0, ref 0)
      a = spawn (lambda () . fib (x - 1) da)
      b = spawn (lambda () . fib (x - 2) db)
      (((),())) = (sync a, sync b)
    in
      dest ← !da + !db
    end
```

```
fib x =
  if x ≤ 1 then
    x
  else
    let (ra, rb) = (fib (x - 1)) || (fib (x - 2)) in
      ra + rb
    end.
```



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Mutable State and Race Conditions

- ***Parallel* versus *Sequential Semantics***

- We can think of any parallel SPARC program as a sequential program by replacing parallel tuples with sequential ones
- ***Sequential Elision***
 - obtained by replacing `par` (or `||`) with simple sequential pairs
- if the original parallel program is pure (purely functional), and does not use side effects, then corresponding sequential program is “*observationally equivalent*” to the parallel one



Mutable State

- a multithreaded or parallel program uses *mutable state*
 - reasoning about its correctness and efficiency can become difficult

```
let x = ref 0
    (((),()) = (x ← 1) || (x ← 2)
in print x end
```

```
let x = ref 0
    (((),()) = (x ← x + 1) || (x ← x + 1)
in print x end.
```



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Data race

- a data race occurs
 - multiple threads access **the same piece of data** and **at least one of the threads update or write** to the data
 - usually lead to ***non-deterministic outcome***, which is in many cases an error condition
 - ***benign*** : a data race does not impact adversely the correctness of a program



Why Use Mutable State

- Why use mutable state at all and why not program purely functionally?
 - it is **impossible to avoid mutable state completely**
 - Even if a program is purely functional, it still has to allocate memory and write into it.
 - at some level of abstraction, we must operate on mutable state, even when it is “hidden” behind the abstraction of purely functional programming
 - Mutable state **enables implementing certain operations including purely functional ones more efficiently**
 - updating a single position in an array requires copying an array if we are not allowed to mutate it



Races are Considered Harmful

- data races are harmful and should be avoided to the extent possible
 - have 100 threads each of which execute 10 instructions that can cause races,
 - the total number of interactions we must consider are 10^{100} more than the number of atoms in the universe



华中科技大学
HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Critical Sections and Mutual Exclusion

- a ***critical section*** is a part that cannot be executed by more than one thread at the same time
 - a critical section must be executed in mutual exclusion
 - typically contain code that alters data shared among parallel computations, and could lead to data races if executed concurrently

```
let
    debit bal delta =
        bal ← bal - delta
    credit bal delta =
        bal ← bal + delta
in
    (debit mybal 10) || (credit mybal 10)
end
```



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Data Races and Critical Sections

- a ***data race*** could occur: a critical section is executed concurrently
 - When a data race occurs, the outcome of the program usually depends on the relative timing of events in the execution, and varies from one execution to another
 - It can be extremely difficult to find a data race
 - parallel programs usually exhibit ample non-determinacy in their execution, due to scheduling
 - Heisenbug
 - A data race may lead to an observably incorrect behavior only a tiny fraction of the time, making it extremely difficult to observe and reproduce it



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Mutual Exclusion Problem

- Solving: synchronizing between the threads by using synchronization instructions
 - **Spin locks** allow a thread to “busy wait” until the critical section is “clear” of other threads
 - **Blocking locks** allow a thread to “block” and wait for another thread to exit the critical section. When the critical section is clear, then the blocked thread receives a signal, allowing it to proceed. The term mutex , short for "mutual exclusion" is sometimes used to refer to a blocking lock
 - **Atomic read-modify-write instructions** , can read and modify the contents of a memory location atomically, allowing a thread to operate safely on shared data



Atomic Read-Modify-Write Instructions

- Nonblocking Synchronization
 - atomic read-modify-write instructions are called *nonblocking*
 - Nonblocking operations can be used to implement more complex concurrent *nonblocking data structures*
- *Compare and Swap*
 - an atomic read-modify-write instruction that performs a memory read followed by a memory write atomically on a machine word
- *Fetch and Add*
 - an atomic read-modify-write instruction that atomically updates the contents of a memory location and returns the contents (before the update)



Summary

- Functional Algorithms
- The Lambda Calculus
- The SPARC Language
 - Syntax and Semantics
 - Type System of SPARC
- Example
 - Threads, Concurrency, and Parallelism
 - Critical Sections and Mutual Exclusion



华中科技大学

HUZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY