

```
In [1]: import os
import random
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader, Subset
from torchvision import datasets, transforms
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

SEED = 40
random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed_all(SEED)

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
device
```

```
Out[1]: device(type='cuda')
```

```
In [ ]: # 配置信息
class Cfg:
    data_root = './data'           # MNIST download/cache dir
    train_fraction = 0.10          # use 10% of the train split
    test_fraction = 0.10           # use 10% of the test split
    train_num_pairs = 10000         # number of training pairs
    test_num_pairs = 2000           # number of test pairs
    batch_size = 64
    epochs = 50
    lr = 1e-3
    weight_decay = 1e-4
    num_workers = 2
    save_dir = './output'
    model_name = 'siamese_mnist_same_digit.pth'
    # batch_log = 'batch_log.log'
    epoch_log = 'epoch_log.csv'

os.makedirs(Cfg.save_dir, exist_ok=True)
Cfg.__dict__
```

```
Out[ ]: mappingproxy({'__module__': '__main__',
                      'data_root': './data',
                      'train_fraction': 0.1,
                      'test_fraction': 0.1,
                      'train_num_pairs': 10000,
                      'test_num_pairs': 2000,
                      'batch_size': 64,
                      'epochs': 50,
                      'lr': 0.001,
                      'weight_decay': 0.0001,
                      'num_workers': 2,
                      'save_dir': './output',
                      'model_name': 'siamese_mnist_same_digit.pth',
                      'epoch_log': 'epoch_log.csv',
                      '__dict__': <attribute '__dict__' of 'Cfg' objects>,
                      '__weakref__': <attribute '__weakref__' of 'Cfg' objects>,
                      '__doc__': None})
```

```
In [3]: # Transforms预处理
transform = transforms.Compose([
    transforms.ToTensor(), # 归一化[0,1], shape (1,28,28)
```

```

        transforms.Normalize((0.1307,), (0.3081,)),
    ])

# 加载MNIST数据集
train_dataset = datasets.MNIST(Cfg.data_root, train=True, download=True, transform=transform)
test_dataset = datasets.MNIST(Cfg.data_root, train=False, download=True, transform=transform)
print(f"原始训练集大小: {len(train_dataset)}")
print(f"原始测试集大小: {len(test_dataset)}")

# 抽取10%的数据
def sample_dataset(dataset, sample_ratio=0.1):
    indices = list(range(len(dataset)))
    sampled_indices = random.sample(indices, int(len(dataset) * sample_ratio))
    return Subset(dataset, sampled_indices)
train_subset = sample_dataset(train_dataset, Cfg.train_fraction)
test_subset = sample_dataset(test_dataset, Cfg.test_fraction)
print(f"采样后训练集大小: {len(train_subset)}")
print(f"采样后测试集大小: {len(test_subset)}")

# 创建配对数据集
class MNISTPairsDataset(Dataset):
    def __init__(self, subset, num_pairs):
        self.subset = subset
        self.num_pairs = num_pairs
        self.labels = [label for _, label in subset]
        self.images = [image for image, _ in subset]
        # 生成相同数字的索引映射
        label_to_indices = {}
        for idx, label in enumerate(self.labels):
            if label not in label_to_indices:
                label_to_indices[label] = []
            label_to_indices[label].append(idx)
        # 创建配对
        self.pairs = []
        self.pair_labels = []
        for _ in range(num_pairs // 2):
            # 正样本对
            label = random.choice(list(label_to_indices.keys()))
            if len(label_to_indices[label]) >= 2:
                idx1, idx2 = random.sample(label_to_indices[label], 2)
                self.pairs.append((idx1, idx2))
                self.pair_labels.append(1)
            # 负样本对
            label1, label2 = random.sample(list(label_to_indices.keys()), 2)
            idx1 = random.choice(label_to_indices[label1])
            idx2 = random.choice(label_to_indices[label2])
            self.pairs.append((idx1, idx2))
            self.pair_labels.append(0)

    def __len__(self):
        return len(self.pairs)

    def __getitem__(self, idx):
        idx1, idx2 = self.pairs[idx]
        x1 = self.images[idx1]
        x2 = self.images[idx2]
        y = self.pair_labels[idx]
        return x1, x2, torch.tensor(y, dtype=torch.float32)

# 创建训练和测试配对数据集
train_pairs_dataset = MNISTPairsDataset(train_subset, num_pairs=Cfg.train_num_pairs)
test_pairs_dataset = MNISTPairsDataset(test_subset, num_pairs=Cfg.test_num_pairs)
print(f"训练配对数据集大小: {len(train_pairs_dataset)}")
print(f"测试配对数据集大小: {len(test_pairs_dataset)}")

# 检查正负样本比例

```

```

train_labels = [label for _, _, label in train_pairs_dataset]
test_labels = [label for _, _, label in test_pairs_dataset]
print(f"训练集正样本比例: {sum(train_labels) / len(train_labels):.3f}")
print(f"测试集正样本比例: {sum(test_labels) / len(test_labels):.3f}")

# 数据加载器
train_loader = DataLoader(train_pairs_dataset, batch_size=Cfg.batch_size, shuffle=True, num_workers=4)
test_loader = DataLoader(test_pairs_dataset, batch_size=Cfg.batch_size, shuffle=False, num_workers=4)

```

原始训练集大小: 60000
 原始测试集大小: 10000
 采样后训练集大小: 6000
 采样后测试集大小: 1000
 训练配对数据集大小: 10000
 测试配对数据集大小: 2000
 训练集正样本比例: 0.500
 测试集正样本比例: 0.500

```

In [4]: # 可视化一些配对样本
def show_pairs(ds: MNISTPairsDataset, n=4):
    # 创建三列: 左图、右图和标签
    fig, axes = plt.subplots(nrows=n, ncols=3, figsize=(3, 1*n))
    for i in range(n):
        x1, x2, y = ds[random.randint(0, len(ds)-1)]
        y = int(y.item())
        axes[i,0].imshow(x1.squeeze(0), cmap='gray')
        axes[i,0].axis('off')
        axes[i,1].imshow(x2.squeeze(0), cmap='gray')
        axes[i,1].axis('off')
        axes[i,2].text(0.5, 0.5, f'Label = {y}', fontsize=16, ha='center', va='center')
        axes[i,2].axis('off')
    plt.tight_layout()
    plt.show()

show_pairs(train_pairs_dataset, n=4)

```



```

In [5]: class ConvEncoder(nn.Module):
    def __init__(self, emb_dim=64):
        super().__init__()
        self.features = nn.Sequential(
            nn.Conv2d(1, 16, kernel_size=3, padding=1),    # 1x28x28 -> 16x28x28
            nn.ReLU(inplace=True),

```

```

        nn.MaxPool2d(2),                                # -> 16x14x14
        nn.Conv2d(16, 32, kernel_size=3, padding=1),    # -> 32x14x14
        nn.ReLU(inplace=True),
        nn.MaxPool2d(2),                                # -> 32x7x7
    )
    self.fc = nn.Sequential(
        nn.Flatten(),
        nn.Linear(32*7*7, 128),
        nn.ReLU(inplace=True),
        nn.Dropout(0.2),
        nn.Linear(128, emb_dim),
    )

def forward(self, x):
    x = self.features(x)
    x = self.fc(x)
    return x

class SiameseSameDigit(nn.Module):
    def __init__(self, emb_dim=64):
        super().__init__()
        self.encoder = ConvEncoder(emb_dim=emb_dim)
        # 使用绝对差和元素级乘积组合特征
        self.head = nn.Sequential(
            nn.Linear(emb_dim*2, 32),
            nn.ReLU(inplace=True),
            nn.Linear(32, 1), # logit
            nn.Sigmoid() # 输出概率
        )

    def forward(self, x1, x2):
        z1 = self.encoder(x1)
        z2 = self.encoder(x2)
        feat = torch.cat([torch.abs(z1 - z2), z1 * z2], dim=1)
        logit = self.head(feat)
        return logit.squeeze(1) # (B,)

    def print_model_parameters(self):
        print(f"{'Layer Name':<40} {'Parameter Shape':<30} {'Param Count':<15}")
        print("=" * 85)
        total_params = 0
        for name, param in self.named_parameters():
            param_count = param.numel()
            total_params += param_count
            print(f"{name:<40} {str(list(param.shape)):<30} {param_count:<15}")
        print("=" * 85)
        print(f"Total Parameters: {total_params}")

model = SiameseSameDigit(emb_dim=64).to(device)
model.print_model_parameters()

```

Layer Name	Parameter Shape	Param Count
encoder.features.0.weight	[16, 1, 3, 3]	144
encoder.features.0.bias	[16]	16
encoder.features.3.weight	[32, 16, 3, 3]	4608
encoder.features.3.bias	[32]	32
encoder.fc.1.weight	[128, 1568]	200704
encoder.fc.1.bias	[128]	128
encoder.fc.4.weight	[64, 128]	8192
encoder.fc.4.bias	[64]	64
head.0.weight	[32, 128]	4096
head.0.bias	[32]	32
head.2.weight	[1, 32]	32
head.2.bias	[1]	1

Total Parameters: 218049

```
In [ ]: def evaluate(model, loader, loss_fn):
    model.eval()
    total_loss = 0.0
    total_correct = 0
    total = 0
    with torch.no_grad():
        for x1, x2, y in loader:
            x1, x2, y = x1.to(device), x2.to(device), y.to(device)
            outputs = model(x1, x2)
            loss = loss_fn(outputs, y)
            total_loss += loss.item() * y.size(0)
            preds = (outputs >= 0.5).float()
            total_correct += (preds == y.view_as(preds)).sum().item()
            total += y.size(0)
    return total_loss/total, total_correct/total

def train(model, train_loader, test_loader, epochs, lr, weight_decay):
    # loss_fn = nn.BCEWithLogitsLoss() # 内置sigmoid层
    loss_fn = nn.BCELoss()
    # optimizer = optim.Adam(model.parameters(), lr=lr, weight_decay=weight_decay)
    optimizer = optim.SGD(model.parameters(), lr=lr, momentum=0.9, weight_decay=weight_decay)

    history = {
        'epoch': [],
        'train_loss': [],
        'train_acc': [],
        'test_loss': [],
        'test_acc': []
    }
    # log_file_path = os.path.join(Cfg.save_dir, Cfg.batch_log)

    for epoch in range(1, epochs+1):
        model.train()
        epoch_loss = 0.0
        epoch_correct = 0
        epoch_total = 0

        for b, (x1, x2, y) in enumerate(train_loader):
            x1, x2, y = x1.to(device), x2.to(device), y.to(device)

            optimizer.zero_grad()
            outputs = model(x1, x2)
            loss = loss_fn(outputs, y)
            loss.backward()
            optimizer.step()

            with torch.no_grad():
                preds = (outputs >= 0.5).float()
                correct = (preds == y.view_as(preds)).sum().item()
```

```

# batch_loss = Loss.item()
# batch_acc = correct / y.size(0)
epoch_loss += loss.item() * y.size(0)
epoch_correct += correct
epoch_total += y.size(0)
# test_loss_batch, test_acc_batch = evaluate(model, test_loader, loss_fn)
# with open(log_file_path, 'a') as log_file:
#     log_file.write(f"[epoch {epoch:2d}] [batch_idx {b:3d}] " +
#                 f"batch_loss={batch_loss:.4f} batch_acc={batch_acc:.4f} | " +
#                 f"test_loss={test_loss_batch:.4f} test_acc_batch={test_acc_batch:.4f}\n")
#
train_loss = epoch_loss/epoch_total
train_acc = epoch_correct/epoch_total
test_loss, test_acc = evaluate(model, test_loader, loss_fn)

history['epoch'].append(epoch)
history['train_loss'].append(train_loss)
history['train_acc'].append(train_acc)
history['test_loss'].append(test_loss)
history['test_acc'].append(test_acc)

print(f"[Epoch {epoch:2d}] " +
      f"train_loss={train_loss:.4f} train_acc={train_acc:.4f} | " +
      f"test_loss={test_loss:.4f} test_acc={test_acc:.4f}")

# print(f"Training log (mini-batch) saved to {log_file_path}")
hist_df = pd.DataFrame(history)
csv_path = os.path.join(Cfg.save_dir, Cfg.epoch_log)
hist_df.to_csv(csv_path, index=False)
print(f"Training log (epoch) saved to {csv_path}")
model_path = os.path.join(Cfg.save_dir, Cfg.model_name)
torch.save(model, model_path)
print(f"Model saved to {model_path}")

return history

```

history = train(model, train_loader, test_loader, Cfg.epochs, Cfg.lr, Cfg.weight_decay)

[Epoch 1] train_loss=0.6931	train_acc=0.4975	test_loss=0.6930	test_acc=0.4945
[Epoch 2] train_loss=0.6929	train_acc=0.5157	test_loss=0.6927	test_acc=0.5255
[Epoch 3] train_loss=0.6928	train_acc=0.5255	test_loss=0.6924	test_acc=0.5405
[Epoch 4] train_loss=0.6925	train_acc=0.5363	test_loss=0.6921	test_acc=0.5405
[Epoch 5] train_loss=0.6922	train_acc=0.5733	test_loss=0.6916	test_acc=0.6445
[Epoch 6] train_loss=0.6919	train_acc=0.5534	test_loss=0.6910	test_acc=0.6800
[Epoch 7] train_loss=0.6913	train_acc=0.5761	test_loss=0.6902	test_acc=0.6235
[Epoch 8] train_loss=0.6904	train_acc=0.6291	test_loss=0.6889	test_acc=0.6625
[Epoch 9] train_loss=0.6890	train_acc=0.6300	test_loss=0.6868	test_acc=0.6620
[Epoch 10] train_loss=0.6867	train_acc=0.6157	test_loss=0.6833	test_acc=0.6695
[Epoch 11] train_loss=0.6824	train_acc=0.6236	test_loss=0.6769	test_acc=0.6695
[Epoch 12] train_loss=0.6743	train_acc=0.6402	test_loss=0.6641	test_acc=0.6660
[Epoch 13] train_loss=0.6566	train_acc=0.6519	test_loss=0.6376	test_acc=0.6770
[Epoch 14] train_loss=0.6207	train_acc=0.6914	test_loss=0.5885	test_acc=0.7165
[Epoch 15] train_loss=0.5630	train_acc=0.7286	test_loss=0.5293	test_acc=0.7555
[Epoch 16] train_loss=0.5090	train_acc=0.7593	test_loss=0.4792	test_acc=0.7810
[Epoch 17] train_loss=0.4541	train_acc=0.7922	test_loss=0.4280	test_acc=0.8115
[Epoch 18] train_loss=0.4047	train_acc=0.8237	test_loss=0.3920	test_acc=0.8260
[Epoch 19] train_loss=0.3710	train_acc=0.8380	test_loss=0.3599	test_acc=0.8425
[Epoch 20] train_loss=0.3345	train_acc=0.8599	test_loss=0.3277	test_acc=0.8595
[Epoch 21] train_loss=0.3146	train_acc=0.8681	test_loss=0.3138	test_acc=0.8655
[Epoch 22] train_loss=0.2876	train_acc=0.8821	test_loss=0.2814	test_acc=0.8840
[Epoch 23] train_loss=0.2709	train_acc=0.8889	test_loss=0.2682	test_acc=0.8925
[Epoch 24] train_loss=0.2502	train_acc=0.8988	test_loss=0.2485	test_acc=0.8980
[Epoch 25] train_loss=0.2346	train_acc=0.9072	test_loss=0.2386	test_acc=0.9070
[Epoch 26] train_loss=0.2200	train_acc=0.9144	test_loss=0.2153	test_acc=0.9140
[Epoch 27] train_loss=0.2058	train_acc=0.9191	test_loss=0.2141	test_acc=0.9160
[Epoch 28] train_loss=0.1883	train_acc=0.9277	test_loss=0.1980	test_acc=0.9185
[Epoch 29] train_loss=0.1765	train_acc=0.9338	test_loss=0.1819	test_acc=0.9305
[Epoch 30] train_loss=0.1632	train_acc=0.9395	test_loss=0.1767	test_acc=0.9335
[Epoch 31] train_loss=0.1550	train_acc=0.9429	test_loss=0.1579	test_acc=0.9385
[Epoch 32] train_loss=0.1443	train_acc=0.9431	test_loss=0.1530	test_acc=0.9415
[Epoch 33] train_loss=0.1326	train_acc=0.9512	test_loss=0.1472	test_acc=0.9445
[Epoch 34] train_loss=0.1253	train_acc=0.9533	test_loss=0.1340	test_acc=0.9495
[Epoch 35] train_loss=0.1132	train_acc=0.9552	test_loss=0.1346	test_acc=0.9485
[Epoch 36] train_loss=0.1099	train_acc=0.9594	test_loss=0.1456	test_acc=0.9455
[Epoch 37] train_loss=0.0984	train_acc=0.9635	test_loss=0.1183	test_acc=0.9530
[Epoch 38] train_loss=0.0935	train_acc=0.9647	test_loss=0.1190	test_acc=0.9555
[Epoch 39] train_loss=0.0874	train_acc=0.9689	test_loss=0.1259	test_acc=0.9560
[Epoch 40] train_loss=0.0835	train_acc=0.9700	test_loss=0.1234	test_acc=0.9550
[Epoch 41] train_loss=0.0741	train_acc=0.9734	test_loss=0.1119	test_acc=0.9615
[Epoch 42] train_loss=0.0684	train_acc=0.9774	test_loss=0.1098	test_acc=0.9625
[Epoch 43] train_loss=0.0698	train_acc=0.9742	test_loss=0.1063	test_acc=0.9650
[Epoch 44] train_loss=0.0645	train_acc=0.9780	test_loss=0.1071	test_acc=0.9610
[Epoch 45] train_loss=0.0584	train_acc=0.9779	test_loss=0.1185	test_acc=0.9575
[Epoch 46] train_loss=0.0585	train_acc=0.9792	test_loss=0.1066	test_acc=0.9625
[Epoch 47] train_loss=0.0569	train_acc=0.9789	test_loss=0.0959	test_acc=0.9650
[Epoch 48] train_loss=0.0450	train_acc=0.9856	test_loss=0.1092	test_acc=0.9645
[Epoch 49] train_loss=0.0499	train_acc=0.9821	test_loss=0.0927	test_acc=0.9675
[Epoch 50] train_loss=0.0476	train_acc=0.9828	test_loss=0.1007	test_acc=0.9660

Training log (epoch) saved to ./output/epoch_log.csv

Model saved to ./output/siamese_mnist_same_digit.pth

```
In [7]: def plot_results(history):
    # 绘制损失和准确率曲线
    fig, axes = plt.subplots(1, 2, figsize=(12, 4))

    # 左图: 损失曲线
    axes[0].plot(history['train_loss'], label='Train Loss')
    axes[0].plot(history['test_loss'], label='Test Loss')
    axes[0].set_xlabel('Epoch')
    axes[0].set_ylabel('Loss')
    axes[0].set_title('Loss Curves')
    axes[0].legend()

    # 右图: 准确率曲线
    axes[1].plot(history['train_acc'], label='Train Acc')
    axes[1].plot(history['test_acc'], label='Test Acc')
    axes[1].set_xlabel('Epoch')
    axes[1].set_ylabel('Accuracy')
    axes[1].set_title('Accuracy Curves')
    axes[1].legend()
```

右图: 准确率曲线


```

        axes[1].plot(history['train_acc'], label='Train Acc')
        axes[1].plot(history['test_acc'], label='Test Acc')
        axes[1].set_xlabel('Epoch')
        axes[1].set_ylabel('Accuracy')
        axes[1].set_title('Accuracy Curves')
        axes[1].legend()

    plt.tight_layout()
    plt.show()

plot_results(history)

```



```

In [8]: def plot_confusion_matrix(model, loader):
    model.eval()
    y_true = []
    y_pred = []


    with torch.no_grad():
        for x1, x2, y in loader:
            x1, x2, y = x1.to(device), x2.to(device), y.to(device)
            outputs = model(x1, x2)
            preds = (outputs >= 0.5).float()
            y_true.extend(y.cpu().numpy())
            y_pred.extend(preds.cpu().numpy())

    # 计算混淆矩阵
    cm = confusion_matrix(y_true, y_pred)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Different", "Same"])
    disp.plot()
    plt.title("Confusion Matrix")
    plt.show()

plot_confusion_matrix(model, test_loader)

```

Confusion Matrix




```
In [9]: def predict_random_pair(model, test_subset):
    # 随机挑选两张图片
    idx1, idx2 = random.sample(range(len(test_subset)), 2)
    x1, label1 = test_subset[idx1]
    x2, label2 = test_subset[idx2]
    # 将图片转换为批量形式并移动到设备
    x1 = x1.unsqueeze(0).to(device) # (1, 1, 28, 28)
    x2 = x2.unsqueeze(0).to(device) # (1, 1, 28, 28)
    # 模型预测
    model.eval()
    with torch.no_grad():
        output = model(x1, x2) # 输出 logits
        pred = 1 if output >= 0.5 else 0 # 二分类预测
    # 显示图片和预测结果
    fig, axes = plt.subplots(1, 2, figsize=(6, 3))
    axes[0].imshow(x1.cpu().squeeze(0).squeeze(0), cmap='gray')
    axes[0].set_title(f"Label: {label1}")
    axes[0].axis('off')
    axes[1].imshow(x2.cpu().squeeze(0).squeeze(0), cmap='gray')
    axes[1].set_title(f"Label: {label2}")
    axes[1].axis('off')
    result_text = f"Predicted: {'Same' if pred == 1 else 'Different'} | Actual: {'Same' if la
    fig.suptitle(result_text, fontsize=14, y=0.04)
    plt.tight_layout()
    plt.show()

model = torch.load(os.path.join(Cfg.save_dir, Cfg.model_name), map_location=device, weights_o
model = model.to(device)
predict_random_pair(model, test_subset)
```

Label: 3



Label: 8



Predicted: Different | Actual: Different